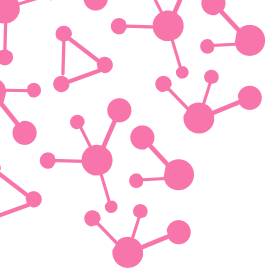

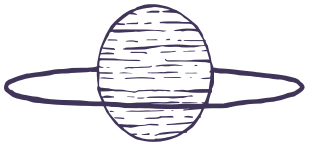

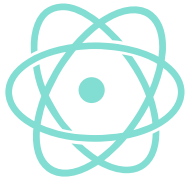

Data Science At AWS

Rui LIU, Bruno BALDEZ CORREA

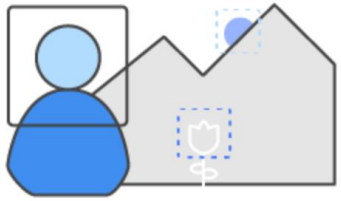

1

Data Science Services at AWS

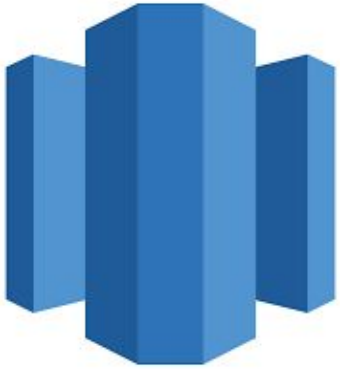

2 Amazon Redshift

Amazon Redshift is a fast, fully managed, petabyte-scale cloud-base data warehouse solution offered by Amazon Web Services that provides simple and cost-effective functionalities to analyze all your data using standard SQL and BI techniques.

Characteristics

- Fast
- Inexpensive
- Extensible
- Simple
- Scalable
- Secure
- Compatible

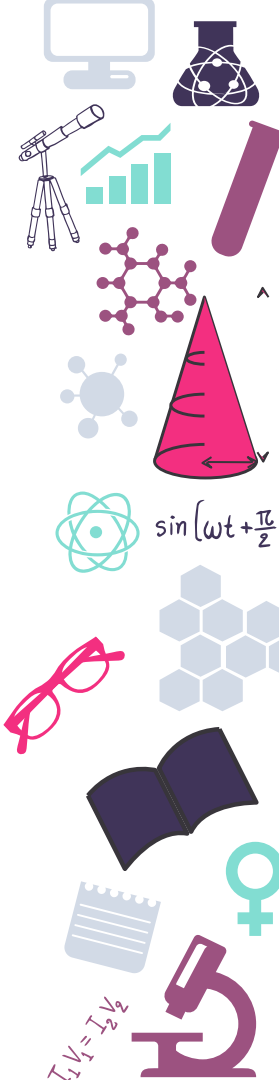

Architecture

Clusters

- 1 or more compute nodes.

Leader node

- compiles code
- distributes it

Compute nodes

- compiles code
- execute the compiled code
- send intermediate results
- dense storage nodes or dense compute nodes

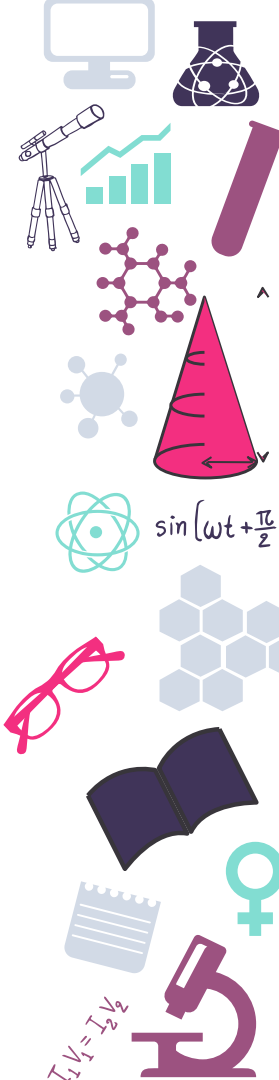

Architecture

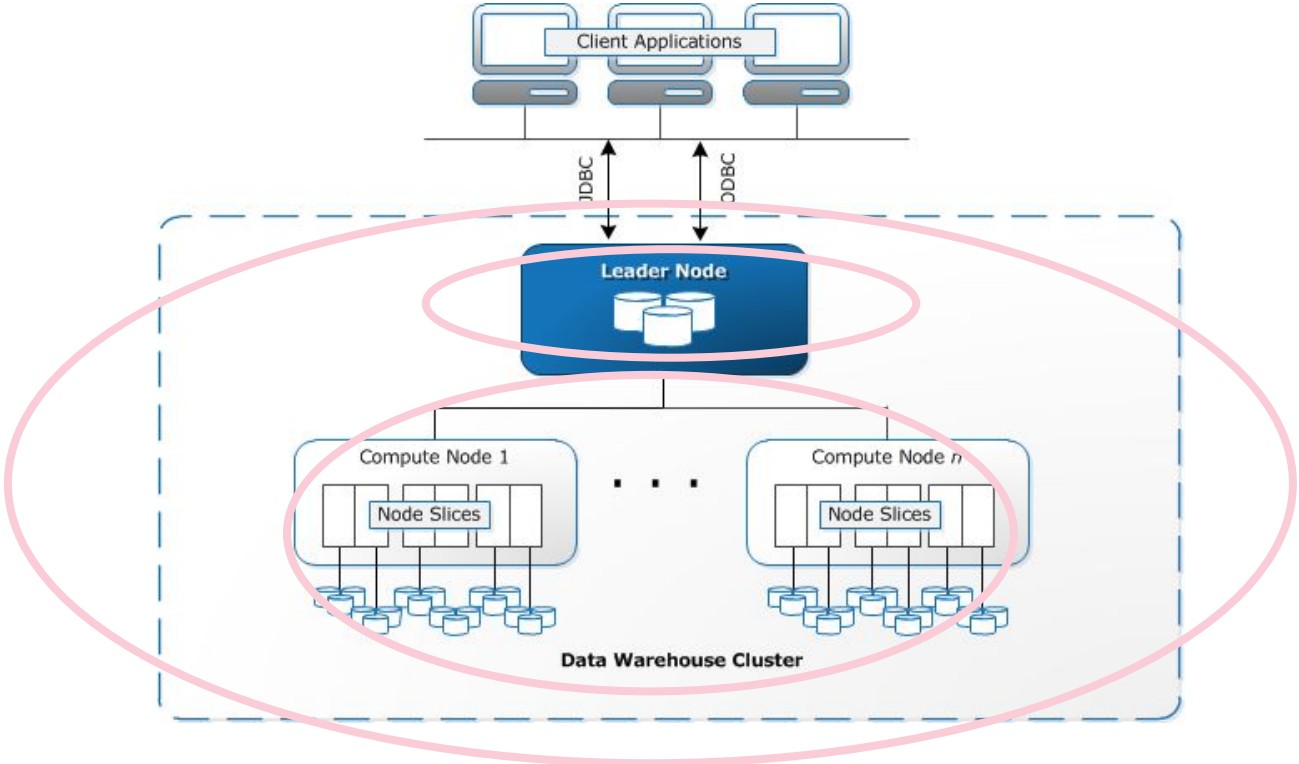

Data Distribution

- ALL
- EVEN
- KEY

- SortKey
- Primary and Foreign keys

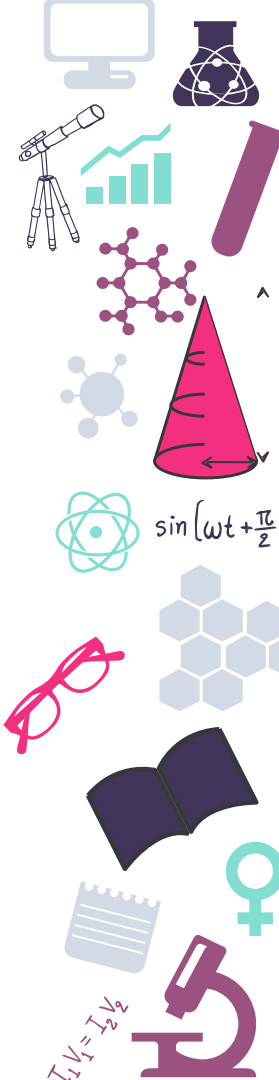

Getting Started with Amazon Redshift

Step 1: Set Up Prerequisites
Step 2: Create an IAM Role
Step 3: Launch a Sample Amazon Redshift Cluster
Step 4: Authorize Access to the Cluster
Step 5: Connect to the Sample Cluster
Step 6: Load Sample Data from Amazon S3
Step 7: Find Additional Resources and Reset Your Environment

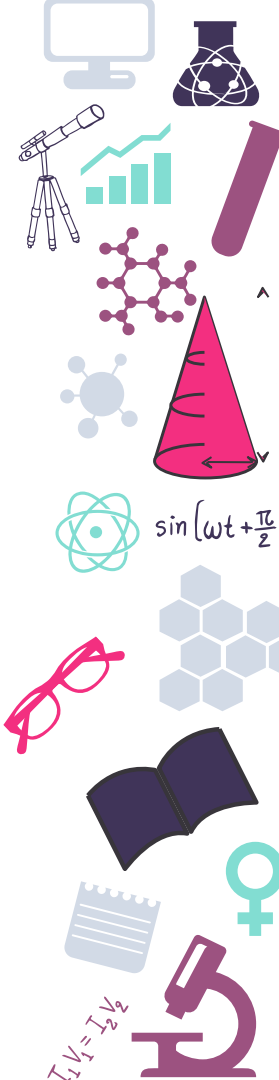

Redshift Use Scenario

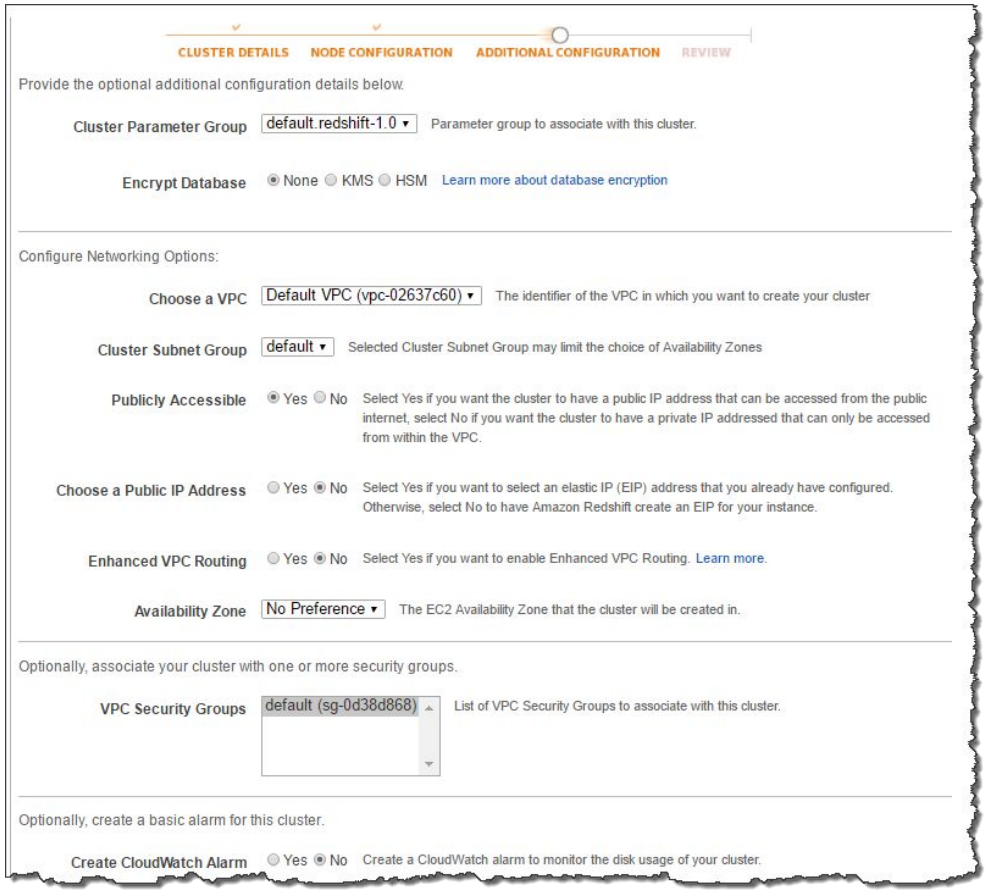

Redshift Use Scenario

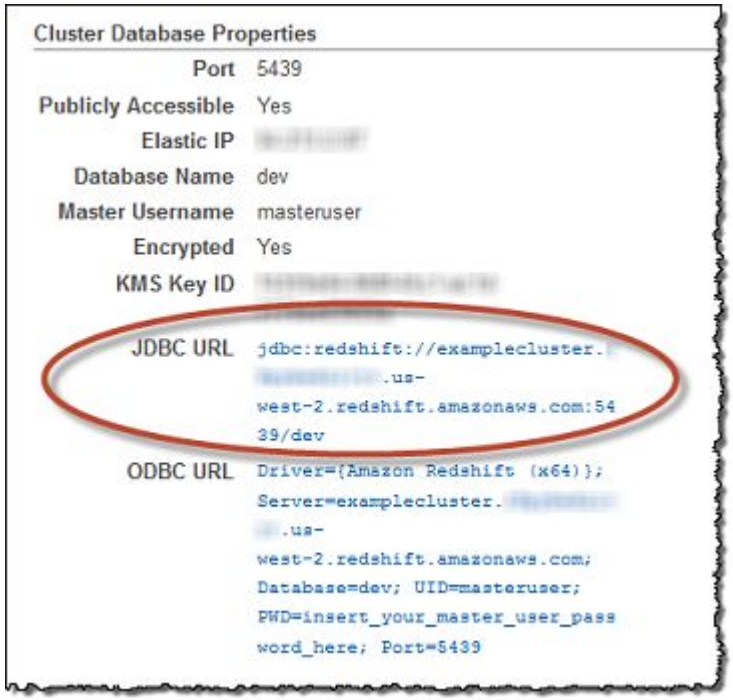

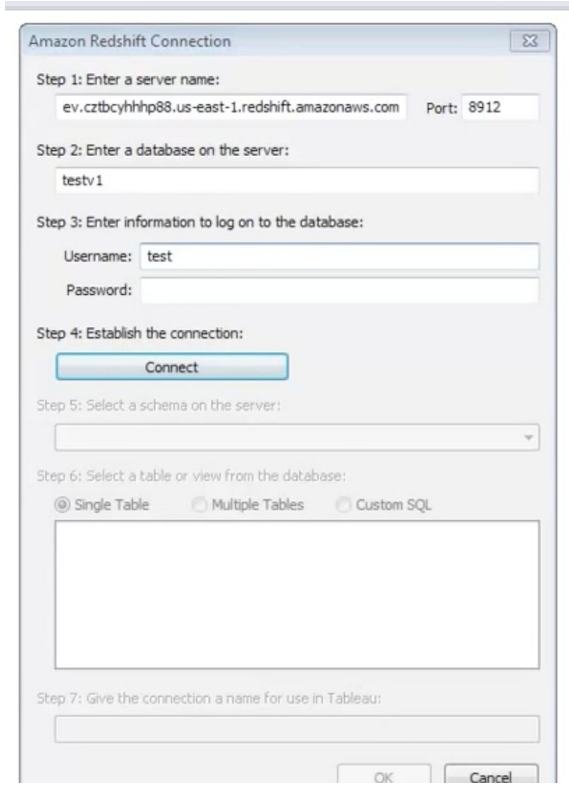

Redshift Use Scenario

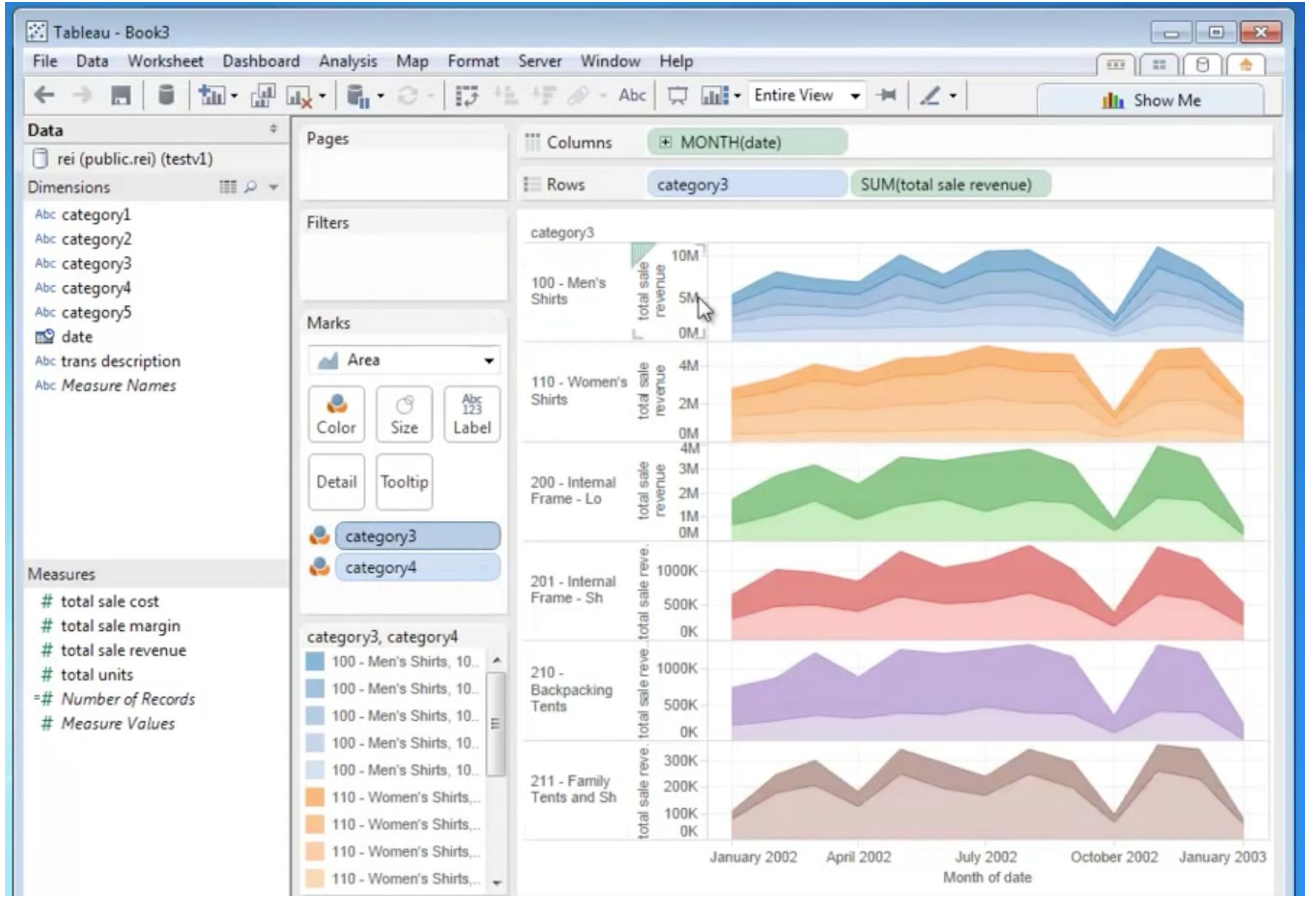

3 Amazon Machine Learning

Amazon Machine Learning makes it easy for developers to build machine learning model without learning complex algorithm or hiring experts.

Ideal Usage Pattern

Hard to code rules

- Rules are not explicit
- Number of factors are huge

Hard to scale

- Large number of tasks
- Impossible to classify tasks manually

Datasources

Datasource is an object used by Amazon Machine Learning as train data, evaluation data and validation data

- Data should be well-formatted
- Datasources should contain one column as target

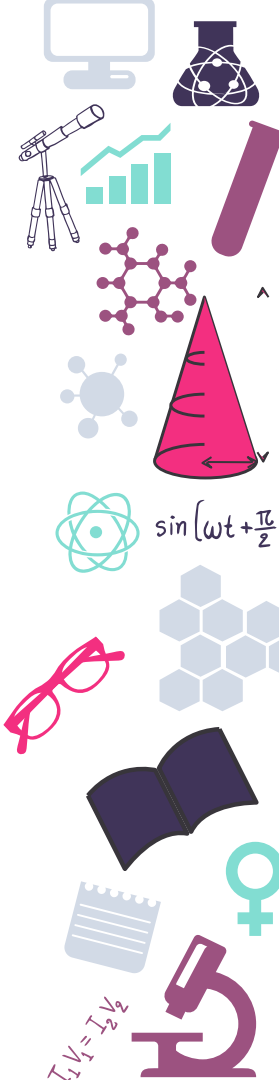

Train ML Models

Amazon ML applies machine learning algorithms automatically

- Binary Classification
- Multiclass Classification
- Regression Model

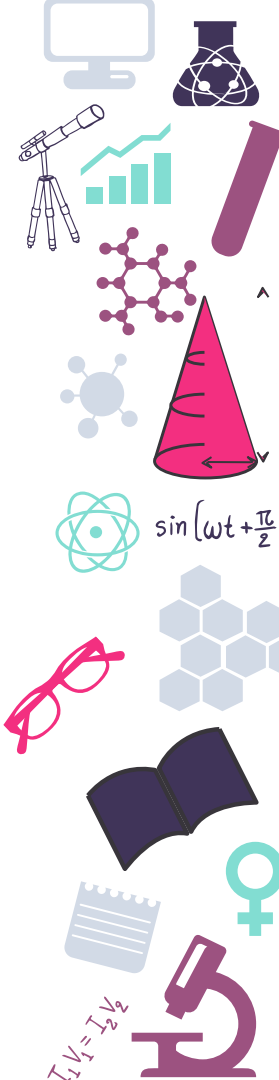

Evaluate ML Models

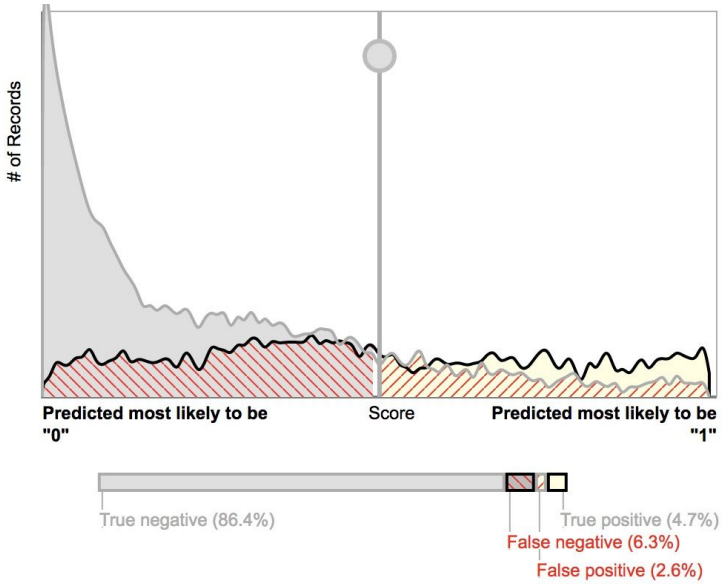

Trade-off based on score threshold 0.5 Reset score threshold (0.5)

- **91% are correct**
  581 true positive
  10,676 true negative
- **9% are errors**
  316 false positive
  782 false negative

- 7% of the records are predicted as "1"
- 93% of the records are predicted as "0"

Save score threshold at 0.50

Advanced metrics

False positive rate **0.0287**	0	1
Precision **0.6477**	0	1
Recall **0.4263**	0	1
Accuracy **0.9111**	0	1

Generate and Interpret Prediction

- Batch Prediction
- Real-Time Prediction

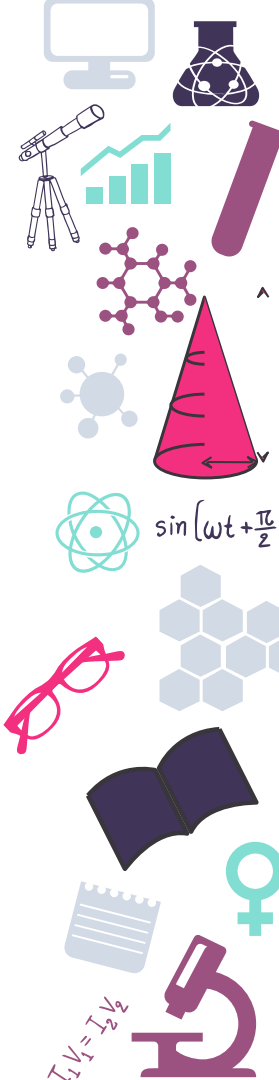

Cost Model

Size of Model

Number of Predictions

Advantages & Disadvantages

Advantages

- Automatic
- Fast and easy

Disadvantages

- Black box
- Supervised model only

Thanks!

Any questions?